# A Comparison of API-based JIT Compiler Overhead during Run-time

Eric Coffin
Faculty of Computer Science
University of New Brunswick
Fredericton, NB, Canada
eric.coffin@unb.ca

## ABSTRACT

Just-in-Time compilation has allowed for significant performance gains during the run-time of applications. Two popular open-source JIT frameworks are LLVM MCJIT and OMR JitBuilder, both of which can be embedded within applications, offering interaces to define and generate native code at run-time. LLVM is a collection of modular compiler and toolchain components, while MCJIT is a framework based on these components to provide JIT compilation. Similarly, JitBuilder is a JIT framework based on the OMR compiler and runtime components. In this report we discuss the different approaches these two frameworks employ. In addition, we measure the overhead of each framework while compiling and then executing a small handful of functions. We found that while LLVM required a larger memory footprint, in certain cases it was able to generate code more quickly. Furthermore, the code LLVM generated typically offered high throughput. JitBuilder, a relatively young project compared to LLVM, does not currently expose the underlying configuration of TRJIT. Instead, it locks the compilation level to the *warm* setting, thus limiting the oppourtunities for optimizations.

## CCS CONCEPTS

• **Software and its engineering → Compilers**.

## KEYWORDS

compilers, just-in-time, optimization

## 1 INTRODUCTION

Just-in-time compilation, or JIT compilation, is a technique to improve the run-time binary translation of an application [1]. Using information collected from the running application, JIT compilers can further optimize generated code. For instance, by collecting profile information on executed code paths, JIT compilers can generate code that is optimized for hot-paths [2]. A JIT compiler might also inline entire blocks of code, or generate an inline cache to speed up the dispatch of polymorphic method calls [3].

Popular high-level language runtimes for Java make heavy use of JIT compilers, allowing their workloads, to execute much faster than if they were entirely interpreted [4]. Given that JIT compilation can add significant overhead to a workload, compilation is typically applied selectively to the code executed most frequently. Furthermore, given that compilation can be viewed as a continuum, where slow, unoptimized code is cheap to generate, and where

fast, optimized code is expensive to generate, a compiler will often support multiple optimization levels [2]. A runtime with such an optimizing JIT compiler can then employ a staged compilation strategy which would allow the runtime to apply compilation and optimization according to heuristics [5].[1]

Considering JIT compilers may also need perform some of the optimizations found in a static compiler such as common sub-expression elimination, loop unrolling and constant propagation [9], an application developer, or language designer interested in enhancing run-time performance by adding a JIT compiler, will have a large engineering task ahead of them. In addition, for projects targeting multiple architectures, considering JIT compilers generate native, or architecture specific code, the effort required will increase significantly. It is no surprise then, that libraries or frameworks, encapsulating proven, high-performance JIT compilers are available to software engineers today. For an engineer interested in incorporating an existing JIT framework, it is useful to consider the following questions when making their selection:

- How much larger is the application binary with the inclusion of the JIT framework?
- How is the resident state set for the application effected?
- How much processor overhead does the JIT compiler have?
- How configurable is the JIT framework?
- What is the quality of the generated JIT code and what effect does this have on throughput?
- How easy is it for the programmer to incorporate and use the JIT framework?

In this report, we will consider those questions while looking at two such JIT frameworks: LLVM MCJIT [10] and OMR JitBuilder [11]. In Sections 2, and 3 we will discuss the background of each compiler and consider the techniques they employ. In Section 3, we outline the methods we used to answer those questions. In Section 4, we will consider the results. In Section 5, we will discuss related work. In Section 6, we will look at potential future work, and finally in Section 7, we will summarize our findings.

## 2 BACKGROUND

In this section we will discuss the background of the two JIT compiler frameworks we are interested in: LLVM MCJIT and OMR JitBuilder. In particular, we will focus on the motivation for each

---

[1] An example of a staged, or tiered compilation strategy can seen with the Testarossa JIT compiler (TRJIT) in the OpenJ9 JVM [6, 7]. Here an invocation threshold must be met before the JIT compiler will compile a method. Additionally, separate thresholds can be associated with each optimization level [8]. We will discuss TRJIT in more detail when we look at JitBuilder in Section 2.2.

```
1  define i32 @mul_add(i32 %x, i32 %y, i32 %z) {
2      entry:
3      %tmp = mul i32 %x, %y
4      %tmp2 = add i32 %tmp, %z
5      ret i32 %tmp2
6  }
```

**Listing 1: LLVM IR for a function multiplying x * y and adding z [13].**

```
1  treetop--> istore a
2                |
3             imul--- isub-------+
4              |              |
5             iadd            |
6              |              |
7             +-----> iload b <--+
8              |              |
9             +-----> iload a <--+
10 }
```

**Listing 2: OMR IR representation for (a+b)*(a-b). Note that the iload nodes are reused [19].**

framework, as well as discuss the techniques and features they provide.

## 2.1 LLVM

LLVM, which at one time stood for Low Level Virtual Machine, is a popular set of open-source, modular compiler and toolchain components [12]. The compiler framework was originally designed to provide analysis and transformation for an application throughout it's entire lifetime: from initial compilation and linking, through to runtime and even while the application was offline (see Figure 1). To achieve this ambitious goal, the framework utilizes a well defined, human-readable, intermediate representation called LLVM IR. The IR, which is initially generated by the front end can be packaged with the target architecture binary along with profiling instructions for later runtime compilation (JIT) as well as more aggressive offline optimizations. Several important characteristics of LLVM IR as as follows:

- The IR maintains Static Single Assignment (SSA) form with unlimited virtual registers.
- Each register is of one of four primitive types: boolean, integer, floating-point or pointer.
- Similar to RISC, memory operations are carried out in registers, and between registers and memory using Load and Store instructions.
- The IR is limited to 31 opcodes.
- The IR is organized into basic blocks which must be composed into valid control flow graphs, simplifying the work required for various optimizations.

This report will focus on LLVM's JIT component, which can be accessed through the MCJIT API. The MCJIT framework provides an API thats accepts IR, generates optimized machine code, and provides a function pointer for calling the generated code. The JIT compiler offers several levels of optimization: none, less, default, and aggressive. It should be noted that the JIT compiler by default does not perform any IR optimizations or transformations. Instead, a developer must pass the generated IR to a PassManager with specific optimizations they intend to apply. These passes can be categorized as analysis passes, or transformation passes [14]:

- Analysis Passes: collect information about IR for use later by transformations, for debugging or for visualization. A few examples are *print-callgraph*, *print-function*, and *iv-users* for printing the users of a particular induction variable. There are roughly 40 such passes available.
- Transformation Passes: These typically modify IR. Examples include *adce* for dead code elimination, *instcombine* for combining redundant instructions, or peephole optimizing, and

*tailcallelim* for eliminating tail calls. There are roughly 60 such passes available.

This optimized IR will then be used by the ExecutionEngine when generating code for the target architecture. Before a function is executed by the ExecutionEngine, it first checks if the ObjectCache contains a copy. If the function could not be found, the compiler will generate code and store it in the ObjectCache before execution [15]. It is worth noting that an newer JIT API, called ORCJIT, is also part of the LLVM project. ORCJIT, or On-Request-Compilation JIT, is intended to compliment the MCJIT API – which compiles eagerly, by adding support for lazy, and concurrent compilations [16].

## 2.2 JitBuilder

JitBuilder is the embeddable framework for interacting with the Eclipse OMR JIT compiler Testarossa [11]. Eclipse OMR is an open-source collection of components for building language runtime environments. Some of OMR's components include a garbage collection framework, a thread library, cross-platform port support, virtual machine building blocks, and a JIT compiler [6, 17]. Much of the infrastructure driving Eclipse OpenJ9, a popular, open-source Java Virtual Machine, links to OMR components.

Through JitBuilder's API, users generate IR upon which optimizations are applied and from which native code is generated [18]. Based on basic-blocks, the IR is arranged into directed acyclic graphic (DAG) structures called trees, composed of nodes which each contain an opcode. OMR IR has several hundred opcodes, where typically with each type operation has a code for each data type: integer, pointer, double, float, vector, etc...[2] The children of these opcode nodes are in turn operands. Trees with side-effects, which cannot be reordered, belong to a list of elements called treetops [19]. Existing nodes can be reused within a given tree, making optimizations such as common subexpression elimination relatively simple (see Listing 2). Though Testarossa supports several levels of optimization ranging from *cold*, with roughly 20 optimizations applied, all the way to *scorching*, where as many as 170 optimizations are applied, as of writing, JitBuilder is locked at the *Warm* optimization level [20, 21].

## 3 QUESTIONS AND METHODS

To answer the questions we outlined in the introduction, we wrote three simple programs for each of LLVM MCJIT, JitBuilder, and a

---

[2] This can be contrasted with LLVM's roughly 31 opcodes, where the data type code is instead embedded within the instruction.

**Figure 1: The LLVM Compiler Framework Architecture [12].**

```
1   bool
2   IncrementMethod::buildIL()
3   {
4     Return(
5       Add(
6           Load("value"),
7           ConstInt32(1)));
8     return true;
9   }
10  }
```

**Listing 3: Generating JitBuilder IR for the increment program.**

```
1   static Function *CreateIncrementFunction(
2                   Module *M,
3                   LLVMContext &c) {
4     FunctionType *f = FunctionType::get(Type::getInt32Ty(c),
5                       {Type::getInt32Ty(c)},
6                       false);
7
8     Function *incrementF = Function::Create(f,
9                       Function::ExternalLinkage,
10                      "increment",
11                      M);
12
13    BasicBlock *BB = BasicBlock::Create(c,
14                      "EntryBlock",
15                      incrementF);
16
17    Value *One = ConstantInt::get(Type::getInt32Ty(c), 1);
18
19    Argument *ArgX = &*incrementF->arg_begin();
20    ArgX->setName("AnArg");
21
22    Value *Sum = BinaryOperator::CreateAdd(ArgX,
23                           One,
24                           "addresult",
25                           BB);
26
27    ReturnInst::Create(c, Sum, BB);
28    return incrementF;
29  }
30  }
```

**Listing 4: Generating MCJIT IR for the increment program.**

straight forward native implementation without a JIT framework. The programs are as follows:

- increment - Calls a single function that adds one to an integer argument.
- recursive-fib - Calls a recursive fibonacci function for *fib(20)*.
- iterative-fib - Calls an iterative fibonacci function for *fib(20)*.

For the JIT implementations, the functions were created using the provided API's, then native code was generated and later used for the actual function call. We built and benchmarked the programs on an x86-64 Linux workstation running Ubuntu 19.04 with 32GB of RAM and an Intel i7-8700 (6-core, 12 thread) processor. Both LLVM, and OMR were built from source from their respective Github repositories [22, 23]. All the source code for this project can be found online through Github[24].

## 4 RESULTS

### 4.1 Compilation Time

The first question we looked at was how long each framework took to compile a function 25 times. This time includes running the program from start to finish and thus includes the setup and teardown time of the JIT.

In Table 1 we see that for the increment task JitBuilder's compiled the function 61.5% faster than MCJIT. For the other two tasks, we see that MCJIT was able to compile the functions more quickly than JitBuilder (recursive-fib compiled 16.5% faster, and iterative-fib compiled 40.1% faster ). Looking at the JitBuilder code to generate the function (see Listing **??**), we see a much smaller function body compared to the function used to generate IR for LLVM (see Listing 4). There is no clear reason why this disparity exists. One idea is JitBuilder has a lower base overhead than LLVM, but the work required to generate code for more complex IR favours LLVM.

### 4.2 Execution Time

Once compilation has completed we turn to measuring how much CPU time is spent actually executing the function. To measure this we compiled the function once and then executed the compiled function 1000 times. For each program, the process was repeated 20 times (see Table 2). Aside from the increment program, which favoured JitBuilder, both the recursive and iterative fib functions executed more quickly with LLVM generated code (0.53% faster for recursive, and 0.37% faster for iterative). Compiling the test programs GCC 5.3 with an optimization level of O3 resulted in extremely low execution times for the native increment program as our function calls were optimized away. It is worth noting that the code generated by the LLVM recursive-fib outperformed the native optimized code. Given two of our programs performed more quickly with MCJIT than with JitBuilder, it could be interesting to examine the disassembly of the generated code.[3]

### 4.3 Binary File Size

Looking at the size of the generated files (see Table 3), the linked LLVM MCJIT programs are roughly six times larger in size than the

---

[3]We used GDB to print instructions starting at the function pointer address. The disassembly is using Intel style.

| | LLVM MCJIT | | | Eclipse OMR JitBuilder | | |
|---|---|---|---|---|---|---|
| **Program** | mean (ns) | median | std. dev. | mean (ns) | median | std. dev. |
| increment | 1,378,060 | 1,387,338 | 39,774 | 533,581 | 531,840 | 4911 |
| recursive-fib | 1,548,316 | 1,547,491 | 3925 | 1,854,393 | 1,850,322 | 21,754 |
| iterative-fib | 2,509,502 | 2,519,808 | 60,429 | 4,192,213 | 4,191,730 | 10,419 |

**Table 1: Results of compiling each function 25 times with each JIT framework.**

| | LLVM MCJIT | | | Eclipse OMR JitBuilder | | | Native (C++) | | |
|---|---|---|---|---|---|---|---|---|---|
| **Program** | mean (ns) | median | std. dev. | mean (ns) | median | std. dev. | mean (ns) | median | std. dev. |
| increment | 1,463,310 | 1,463,032 | 1660 | 524,723 | 524,763 | 718 | 0.876 | 0.875 | 0.002 |
| recursive-fib | 13,674,832 | 13,664,199 | 26,548 | 28,707,381 | 28,699,013 | 29,480 | 16,016,325 | 15,996,370 | 60,459 |
| iterative-fib | 2,657,528 | 2,668,792 | 83,388 | 4,227,967 | 4,190,232 | 102,744 | 0.876 | 0.875 | 0.003 |

**Table 2: Results of JIT compiling each function and executing the generated code 1000 times.**

```
1   (gdb) x/24i $1
2     0x7ffff7fcd000 <fib>:           push    rbp
3     0x7ffff7fcd001 <fib+1>:         push    r14
4     0x7ffff7fcd003 <fib+3>:         push    rbx
5     0x7ffff7fcd004 <fib+4>:         cmp     edi,0x2
6     0x7ffff7fcd007 <fib+7>:         jg      0x7ffff7fcd013 <fib+19>
7     0x7ffff7fcd009 <fib+9>:         mov     eax,0x1
8     0x7ffff7fcd00e <fib+14>:        pop     rbx
9     0x7ffff7fcd00f <fib+15>:        pop     r14
10    0x7ffff7fcd011 <fib+17>:        pop     rbp
11    0x7ffff7fcd012 <fib+18>:        ret
12    0x7ffff7fcd013 <fib+19>:        mov     ebx,edi
13    0x7ffff7fcd015 <fib+21>:        lea     edi,[rbx-0x1]
14    0x7ffff7fcd018 <fib+24>:        movabs  r14,0x7ffff7fcd000
15    0x7ffff7fcd022 <fib+34>:        call    r14
16    0x7ffff7fcd025 <fib+37>:        mov     ebp,eax
17    0x7ffff7fcd027 <fib+39>:        add     ebx,0xfffffffe
18    0x7ffff7fcd02a <fib+42>:        mov     edi,ebx
19    0x7ffff7fcd02c <fib+44>:        call    r14
20    0x7ffff7fcd02f <fib+47>:        add     eax,ebp
21    0x7ffff7fcd031 <fib+49>:        pop     rbx
22    0x7ffff7fcd032 <fib+50>:        pop     r14
23    0x7ffff7fcd034 <fib+52>:        pop     rbp
24    0x7ffff7fcd035 <fib+53>:        ret
```

**Listing 5: LLVM MCJIT recursive-fib disassembly (Optimization level 3)**

```
1   (gdb) x/32i $1
2     0x7ffff6a81034: sub     rsp,0x18
3     0x7ffff6a81038: mov     QWORD PTR [rsp+0x10],rbx
4     0x7ffff6a8103d: cmp     edi,0x2
5     0x7ffff6a81040: jl      0x7ffff6a8106d
6     0x7ffff6a81042: mov     rbx,rdi
7     0x7ffff6a81045: lea     edi,[rbx-0x2]
8     0x7ffff6a81048: call    0x7ffff6a81034
9     0x7ffff6a8104d: mov     rdi,rbx
10    0x7ffff6a81050: mov     rbx,rax
11    0x7ffff6a81053: sub     edi,0x1
12    0x7ffff6a81056: call    0x7ffff6a81034
13    0x7ffff6a8105b: xchg    rbx,rax
14    0x7ffff6a8105e: mov     rcx,rbx
15    0x7ffff6a81061: add     eax,ecx
16    0x7ffff6a81063: mov     rbx,QWORD PTR [rsp+0x10]
17    0x7ffff6a81068: add     rsp,0x18
18    0x7ffff6a8106c: ret
19    0x7ffff6a8106d: mov     rax,rdi
20    0x7ffff6a81070: mov     rbx,QWORD PTR [rsp+0x10]
21    0x7ffff6a81075: add     rsp,0x18
22    0x7ffff6a81079: ret
```

**Listing 6: JitBuilder recursive-fib disassembly (Optimization level *Warm*)**

```
1   (gdb) x/2i $1
2     0x7ffff7fcd000 <increment>:        lea     eax,[rdi+0x1]
3     0x7ffff7fcd003 <increment+3>:      ret
```

**Listing 7: llvm increment disassembly (Optimization level 3)**

```
1   (gdb) x/10i 0x7ffff6a81034
2     0x7ffff6a81034:    mov     rax,rdi
3     0x7ffff6a81037:    add     eax,0x1
4     0x7ffff6a8103a:    ret
```

**Listing 8: jitbuilder increment disassembly (Optimization level *Warm*)**

JitBuilder programs. Note that the JIT frameworks were compiled without debug symbols. Part of the explanation for this is may have to do with over-linking the LLVM programs. While there are over 167 possible linkable modules when using LLVM, I found the minimal set required to use MCJIT was just 8 : *core, executionengine, interpreter, passes, mc, mcjit, support, nativecodegen*.

## 4.4 Resident State Set

The resident state (RSS) set is amount of data loaded into RAM during execution at any one time. To measure the maximum RSS, we used the time utility (/usr/bin/time -v) to collect this. Looking at the results in Table 4, we see that LLVM consistently had the largest RSS. The memory overhead required by LLVM is roughly 4.5 times larger than JitBuilder's requirement. Looking at the number of times inactive pages were reclaimed during execution (minor page faults), we see that LLVM had roughly 4 times the number of inactive pages collected.

## 4.5 Developer Experience

One additional aspect to overhead worth considering is the usability of the frameworks from a developer's perspective. Under this banner, we looked at two items: ease of use, and how configurable the frameworks were.

*4.5.1 Usability.* While ease of use is not likely to be the highest priority when considering which JIT framework to select, research has shown that improved API usability can positively impact developer

| Program | LLVM MCJIT | Eclipse OMR JitBuilder | Native (C++) |
|---|---|---|---|
| increment | 60,640,216 | 10,148,608 | 16,616 |
| recursive-fib | 60,671,176 | 10,149,272 | 16,600 |
| iterative-fib | 60,652,728 | 10,149,184 | 16,600 |

Table 3: Total size in bytes of linked binary test programs.

| | LLVM MCJIT | | Eclipse OMR JitBuilder | | Native (C++) | |
|---|---|---|---|---|---|---|
| Program | max RSS (kb) | minor page faults | max RSS (kb) | minor page faults | max RSS (kb) | minor page faults |
| increment | 43,544 | 1478 | 9416 | 344 | 1548 | 65 |
| recursive-fib | 45,472 | 1545 | 9768 | 384 | 1532 | 63 |
| iterative-fib | 44,452 | 1520 | 10,044 | 423 | 1684 | 66 |

Table 4: Maximum resident state set (RSS) in kilobytes during execution of a single compilation and execution of the generated function. Note that Native had no run-time compilation phase. A minor page fault occurs when the OS reclaims inactive pages during execution. Page size was 4096 bytes.

productivity [25], while improving code quality and reducing errors. Overall we found that JitBuilder took less time to integrate, and provided a more intuitive API for generating IR. When considering usability we took note of the following items:

- Both of the frameworks provide public Github repositories [26, 27].
- The instructions for building the frameworks were readily available and simple to follow.
- Linking LLVM took considerable effort as it was unclear which of the 167 objects were required to use MCJIT. After some trial and error, we discovered LLVM provides a utility, llvmconfig, which simplifies this task by generating header and linking arguments for G++ based on a list of objects[4]
- JitBuilder on the other hand required was simple to integrate, requiring only a single header file and single module to link: jitbuilder. On the other hand, LLVM required us to include 20 header files in our programs.
- We found that writing the IR generators in our test programs took less time with JitBuilder than with LLVM. The most challenging task was writing iterative-fib for LLVM, as this required us to grasp the underlying IR requirements which included: an induction value, a Phi node, several basic blocks for before, during and after the loop, condition values, as well as allocation and store pointers[5].
- The in memory objects generated by LLVM MCJIT can contain debug symbols (DWARF format) allowing debugging to take place in GDB 7.0 and higher [28].

*4.5.2 Configuration.* While the API of LLVM MCJIT may have an overabundance of controls, it does provide an high level of configuration compared to JitBuilder, including which optimization and analysis passes to run, control over the memory structure of the generated code, control over the object cache, and what level of optimization to generate the code at. As mentioned, JitBuilder on

the other hand is currently locked at the *Warm* optimization level, limiting the number of optimizations performed on the IR.

## 4.6 Considerations

Considering the JIT compiler behind JitBuilder was designed specifically for dynamic runtime compilation, it makes sense that it has a lighter footprint than MCJIT, which shares it's codebase with the static compiler for LLVM Unlike LLVM's IR, which is meant for life-long usage, the design of OMR's IR does not have ahead of time, or offline goals in mind, thus possibly simplifying the work required whenever IR interactions are made in JitBuilder.

## 5 RELATED WORK

- LLVM IR -> JitBuilder

## 6 FUTURE WORK

Future Work

## 7 SUMMARY

## 8 ACKNOWLEDGEMENTS

## REFERENCES

[1] John Aycock. A Brief History of Just-in-time. *ACM Comput. Surv.*, 35(2):97–113, jun 2003.
[2] Jim Smith and Ravi Nair. *Virtual Machines: Versatile Platforms for Systems and Processes (The Morgan Kaufmann Series in Computer Architecture and Design)*. Morgan Kaufmann Publishers Inc., San Francisco, CA, USA, 2005.
[3] Urs Hölzle, Craig Chambers, and David Ungar. Optimizing dynamically-typed object-oriented languages with polymorphic inline caches. In *Proceedings of the European Conference on Object-Oriented Programming*, ECOOP '91, page 21–38, Berlin, Heidelberg, 1991. Springer-Verlag.

---

[4]Example commands for building the programs can be found makefiles in the accompanying source code [24] under src/<framework>/.
[5]Source code for the LLVM iterative-fibonacci program can be found in the CreateFib-Function within src/llvm/iterative-fib/fib.cpp

[4] Iffat H. Kazi, Howard H. Chen, Berdenia Stanley, and David J. Lilja. Techniques for Obtaining High Performance in Java Programs . *ACM Comput. Surv.*, 32(3):213–240, September 2000.

[5] Mark Leone and R. Dybvig. Dynamo: A staged compiler architecture for dynamic program optimization. 04 1999.

[6] Eclipse Foundation . Eclipse OMR , 2018. https://projects.eclipse.org/projects/technology.omr, Last accessed on 2019-06-19.

[7] IBM Knowledge Center. How the JIT compiler optimizes code . https://www.ibm.com/support/knowledgecenter/en/SSB23S_1.1.0.15/com.ibm.java.vm.80.doc/docs/jit_optimize.html, Last accessed on 2019-06-19.

[8] IBM Knowledge Center. Selectively disabling the JIT or AOT compiler. https://www.ibm.com/support/knowledgecenter/en/SSYKE2_7.0.0/com.ibm.java.zos.70.doc/diag/tools/jitpd_disable_some.html, Last accessed on 2019-06-19.

[9] Keith Cooper and Linda Torczon. *Engineering a Compiler: International Student Edition.* Morgan Kaufmann Publishers Inc., San Francisco, CA, USA, 2003.

[10] LLVM Project . LLVM Compiler Infrastructure . https://llvm.org/, Last accessed on 2020-01-06.

[11] Daryl Maier and Xiaoli Liang. Supercharge a language runtime! In *Proceedings of the 27th Annual International Conference on Computer Science and Software Engineering*, CASCON '17, page 314, USA, 2017. IBM Corp.

[12] Chris Lattner and Vikram Adve. Llvm: A compilation framework for lifelong program analysis & transformation. In *Proceedings of the international symposium on Code generation and optimization: feedback-directed and runtime optimization*, page 75. IEEE Computer Society, 2004.

[13] LLVM Project . A First Function . http://releases.llvm.org/2.6/docs/tutorial/JITTutorial1.html, Last accessed on 2020-01-06.

[14] LLVM Project . LLVM's Analysis and Transform Passes . https://llvm.org/docs/Passes.html, Last accessed on 2020-01-06.

[15] LLVM Project . MCJIT Design and Implementation . https://llvm.org/docs/MCJITDesignAndImplementation.html, Last accessed on 2020-01-06.

[16] LLVM Project . ORC Design and Implementation . https://llvm.org/docs/ORCv2.html, Last accessed on 2020-01-06.

[17] Matthew Gaudet and Mark Stoodley. Rebuilding an airliner in flight: a retrospective on refactoring IBM testarossa production compiler for Eclipse OMR . pages 24–27, 10 2016.

[18] T. Suganuma, T. Ogasawara, M. Takeuchi, T. Yasue, M. Kawahito, K. Ishizaki, H. Komatsu, and T. Nakatani. Overview of the IBM Java Just-in-Time Compiler . *IBM Systems Journal*, 39(1):175–193, 2000.

[19] Eclipse OMR . Testarossa's Intermediate Language: An Intro to Trees . https://github.com/eclipse/omr/blob/master/doc/compiler/il/IntroToTrees.md, Last accessed on 2020-01-06.

[20] Ricardo Nabinger Sanchez, Jose Nelson Amaral, Duane Szafron, Marius Pirvu, and Mark Stoodley. Using machines to learn method-specific compilation strategies. In *International Symposium on Code Generation and Optimization (CGO 2011)*, pages 257–266. IEEE, 2011.

[21] Mark Stoodley . Github issue: JitBuilder only uses warm optimization level #1187 . https://github.com/eclipse/omr/issues/1187, Last accessed on 2020-01-06.

[22] LLVM Project . [mlir][Linalg] Reimplement and extend getStridesAndOffset . https://github.com/llvm/llvm-project/commit/d67c4cc2eb4ddc450c886598b934c111e721ab0c, Last accessed on 2020-01-06.

[23] Eclipse OMR . Merge pull request #4672 from aviansie-ben/tabortwc-be-fix . https://github.com/eclipse/omr/commit/e6f60a17270955f3312ec85f7c0f53a6e6d01957, Last accessed on 2020-01-06.

[24] Eric Coffin . Report Github Repository . https://github.com/edcoffin/vm-project, Last accessed on 2020-01-06.

[25] Brad A. Myers and Jeffrey Stylos. Improving api usability. *Commun. ACM*, 59(6):62–69, May 2016.

[26] LLVM Project . The LLVM Compiler Infrastructure . https://github.com/llvm/llvm-project, Last accessed on 2020-01-06.

[27] Eclipse OMR . JitBuilder . https://github.com/eclipse/omr/tree/master/jitbuilder/release, Last accessed on 2020-01-08.

[28] LLVM Project . Debugging JIT-ed Code With GDB . https://llvm.org/docs/DebuggingJITedCode.html, Last accessed on 2020-01-06.